

An Introduction to hoardeR

Daniel Fischer and Anu Sironen

2016-08-30

Contents

Installation	1
Using hoardeR for identifying cross-species orthologs of novel candidate genes	2
Typical workflow prior the use of hoardeR	2
Extraction of nucleotide sequence of novel candidate region	2
Sending sequences to NCBI	4
Analyze the Blast Results	5
Visualize the intersected Blast results	7

Installation

The installation of **hoardeR** is straight forward, it is located on Cran and can be installed from within R with the command

```
R> install.packages("hoardeR")
```

There are no special system requirements for the installation. However there are a few package dependencies that have to be met. Normally these dependencies are installed automatically, too. If not, missing Cran packages need then to be installed prior to the installation as above, with

```
R> install.packages("packageName")
```

and missing Bioconductor packages (e.g. the package **Biostrings**) with

```
R> source("https://bioconductor.org/biocLite.R")
R> biocLite("Biostrings")
```

The latest developer version, of **hoardeR** is located on GitHub, here

<https://github.com/fischuu/hoardeR>

and the address on Cran with the latest stable release version is

<https://cran.r-project.org/web/packages/hoardeR/>

A webpage for the latest news on **hoardeR** can be found here:

<http://hoarder.danielfischer.name>

Packages from GitHub can be installed directly in R, using the **devtools** package. To install the latest developer version of **hoardeR** from GitHub, run the following commands

```
R> install.packages("devtools")
R> library("devtools")
R> install_github("fischuu/hoardeR")
```

Once the package is installed, it can be loaded into the namespace

```
R> library(hoardeR)
```

Using hoardeR for identifying cross-species orthologs of novel candidate genes

Typical workflow prior the use of hoardeR

The common application of **hoardeR** is to search cross-species orthologs of unannotated, but active regions in a present organism. For that, typically an RNA-seq experiment has been conducted, the reads have been mapped to a reference genome and gene expressions have been estimated using some annotation.

Further, reads from non-annotated regions have been analysed and a set of novel gene candidate regions has been identified. This can be done either across the whole data set or individually for each sample.

From this analysis the user has either a gtf file with novel loci (e.g. Cufflinks provides this), or then a bed file with the new loci.

The first rows of a typical gtf file would like this:

```
1 Cufflinks exon 242203 242862 . + . gene_id "XLOC_000002"; transcript_id "TCONS_00000002"; exon_number "1"; oId "CUFF.2.1"; class_code "u"; tss_id "TSS2";
1 Cufflinks exon 242203 242646 . + . gene_id "XLOC_000002"; transcript_id "TCONS_00000003"; exon_number "1"; oId "CUFF.1.1"; class_code "u"; tss_id "TSS2";
1 Cufflinks exon 254559 256717 . + . gene_id "XLOC_000002"; transcript_id "TCONS_00000003"; exon_number "2"; oId "CUFF.1.1"; class_code "u"; tss_id "TSS2";
1 Cufflinks exon 254240 256717 . + . gene_id "XLOC_000002"; transcript_id "TCONS_00000004"; exon_number "1"; oId "CUFF.3.1"; class_code "u"; tss_id "TSS3";
1 Cufflinks exon 341982 343630 . + . gene_id "XLOC_000003"; transcript_id "TCONS_00000005"; exon_number "1"; oId "CUFF.10.1"; class_code "u"; tss_id "TSS4";
1 Cufflinks exon 342113 342607 . + . gene_id "XLOC_000003"; transcript_id "TCONS_00000006"; exon_number "1"; oId "CUFF.11.1"; class_code "u"; tss_id "TSS5";
1 Cufflinks exon 342961 343494 . + . gene_id "XLOC_000003"; transcript_id "TCONS_00000006"; exon_number "2"; oId "CUFF.11.1"; class_code "u"; tss_id "TSS5";
1 Cufflinks exon 3312599 3313720 . + . gene_id "XLOC_000024"; transcript_id "TCONS_00000073"; exon_number "1"; oId "CUFF.75.1"; class_code "u"; tss_id "TSS37";
1 Cufflinks exon 3446776 3447142 . + . gene_id "XLOC_000024"; transcript_id "TCONS_00000073"; exon_number "2"; oId "CUFF.75.1"; class_code "u"; tss_id "TSS37";
1 Cufflinks exon 9347375 9347527 . + . gene_id "XLOC_000047"; transcript_id "TCONS_00000118"; exon_number "1"; oId "CUFF.115.1"; class_code "u"; tss_id "TSS63";
```

The first lines of a bed file containing the same information would look like this:

```
1 242203 242862 XLOC_000002.1
1 242203 242646 XLOC_000002.2
1 254559 256717 XLOC_000002.3
1 254240 256717 XLOC_000002.4
1 341982 343630 XLOC_000003.1
1 342113 342607 XLOC_000003.2
1 342961 343494 XLOC_000003.3
1 3312599 3313720 XLOC_000024.1
1 3446776 3447142 XLOC_000024.2
1 9347375 9347527 XLOC_000047.1
```

For the identification of cross-species orthologs, the exon structure is not the primary interest. Rather, the whole genomic region that hosts a novel candidate gene is used here. Typically, not all novel gene candidates are considered and the list is filtered according to some criteria, e.g. that a certain amount of samples have to have a minimum number of reads in that region.

The final set of regions of interest is then available in bed format, e.g. like this

```
1 242203 256717 XLOC_000002
1 341982 343630 XLOC_000003
1 3312599 3447142 XLOC_000024
1 9347375 9347527 XLOC_000047
```

Extraction of nucleotide sequence of novel candidate region

For any loci of interest the nucleotide sequences have to be extracted. There are two way to do that. The first option is to do it outside of R using bedtools. For that, the bed file has to be saved to the harddrive first. Assuming the **data.frame** that contains the bed information is called **novelBed** and we have a defined system path to the project, called **projFolder**. An example command to export the data frame **novelBed** then is

```
R> projFolder <- "/home/daniel/MyProjects/hoardeR-Example"
R> exportBed(novelBed, file=file.path(projFolder, "novel.bed"))
```

In the console this bed file can then be used to extract the fasta files, using `bedtools` like this

```
bedtools getfasta -fi <input FASTA> -bed novel.bed -fo novel.fa
```

Here, `<input FASTA>` is a fasta file that contains the genome information of the species under investigation. If the resulting `novel.fa` file is empty or some other errors occur, a common source of error is a mislabeling of the chromosomes between the input fasta file and the corresponding bed file (e.g. leading CHR, Chr, etc.). Extracting the sequences in such a way is especially then advisable, when the species of interest is rare or the fasta file is not available from Ensembl in the latest version.

However, if the species and also the genome assembly version is available from Ensembl, the fasta information can be obtained straight with the `hoardeR` function `getFastaFromBed`. The `hoardeR` package is able to directly download the genomes and annotations of the most common species, a list of available combinations can be found in the `species` dataset, that comes with `hoardeR`

```
R> head(species)
```

	Common.name	Scientific.name	Taxon.ID	Ensembl.Assembly	Accession	Variation.database	Regulation.database	Pre.assembly
1	Aardvark (Pre)	Orycteropus afer afer	1230840	-	-	-	-	OryAfe1
2	Alpaca	Vicugna pacos	30538	vicPac1	-	-	-	-
3	Amazon molly	Poecilia formosa	48698	Poecilia_formosa-5.1.2	GCA_000485575.1	-	-	-
4	Anole lizard	Anolis carolinensis	28377	AnoCar2.0	GCA_000090745.1	-	-	-
5	Armadillo	Dasyurus novemcinctus	9361	Dasnov3.0	GCA_000208655.2	-	-	-
6	Budgerigar (Pre)	Melopsittacus undulatus	13146	-	-	-	-	MelUnd6.3

Of particular interest are here the columns `Scientific.name` and `Ensembl.Assembly`. If your species of interest and your used assembly matches with the information in the `species` table, you can use the automatic `hoardeR` function `getFastaFromBed` with the default settings to obtain your fasta object. Here, we assume that our species of interest is cow respective `bos taurus` and the fasta files should be downloaded to the folder `/home/daniel/fastas/`. From the `species` table we see that, if no further information are provided then the assembly version UMD3.1 is downloaded.

```
R> novelBed <- data.frame(Chr=c(11,18,3),
                          Start=c(72554673, 62550696, 18148822),
                          End=c(72555273, 62551296, 18149422),
                          Gene=c("LOC1", "LOC2", "LOC3"))
```

```
R> myFasta <- getFastaFromBed(novelBed, species="Bos taurus",
+                             fastaFolder="/home/daniel/fastas/")
```

With the same command it is also possible to obtain the genomic information from assemblies that are published under a different Ensembl release, or that are not at Ensembl or that have a newer/older assembly version. For that, the full syntax is

```
R> myFasta <- getFastaFromBed(novelBed, species="Bos taurus", release = "84",
+                             fastaFolder=NULL, version=NULL)
```

Here, the Ensembl release version and also the assembly version can be specified. However, if no release number is given (i.e. `release=NULL`), the function assumes to find a non-Ensembl fasta file in the folder `fastaFolder`. The naming of these fasta files need then to be of the following format, e.g. for the cow assembly version UMD3.1

```
Bos_taurus.UMD3.1.dna.chromosome.1.fa.gz
Bos_taurus.UMD3.1.dna.chromosome.2.fa.gz
Bos_taurus.UMD3.1.dna.chromosome.3.fa.gz
...
```

That means, first the scientific name, with underscores, then a dot, then the assembly identifier followed by another dot, then `dna.chromosome.` and the chromosome identifier, and then the file extension `fa.gz`. That

means, the fasta files have to be gzipped.

Consider now an example where the genome is not available from Ensembl. On Ensembl is only the sheep genome assembly 'Oar_v3.1', if one needs to work with the 'Oar_v4.0', one would download the chromosome-wise files and then store them in some folder on the harddrive as

```
Ovis_aries.Oar_v4.0.dna.chromosome.1.fa.gz
Ovis_aries.Oar_v4.0.dna.chromosome.2.fa.gz
Ovis_aries.Oar_v4.0.dna.chromosome.3.fa.gz
...
```

To obtain then the corresponding fasta objects (of class `fa`), the command is

```
R> myFasta <- getFastaFromBed(novelBed, species="Ovis aries", release = NULL,
                             fastaFolder="/home/daniel/fasta/", version="Oar_v4.0")
```

assuming again that the fasta files are located in `/home/daniel/fasta/`.

The option to extract the fasta sequences via bedtools is, however, much faster and is the recommended way. The internal function is especially useful for small bed files or fast testing purposes. An `fa`-object can be stored to the harddrive via

```
R> exportFA(myFasta, file="/home/daniel/myFasta.fa")
```

A previously stored or downloaded fasta file can be imported to `hoardeR` with the `importFA()` function:

```
R> novelFA <- importFA(file="/home/daniel/myFasta.fa")
```

The names of the fasta sequences should follow the form like

```
>Chr:Start-End
```

That means, for a sequence from chromosome 12 that starts at 123 and ends at 456 the fasta sequence should be named like the following.

```
>12:123-456
```

Sending sequences to NCBI

Once the fasta object is available in R, it can be send to the NCBI blast service, using the central `hoardeR` function `blastSeq()`

```
R> blastSeq(novelFA,
+           email="daniel.fischer@luke.fi",
+           xmlFolder=file.path(projFolder,"hoardeROut/"),
+           logFolder=file.path(projFolder,"hoardeRLog/"),
+           keepInMemory=FALSE)
```

The main parameters are the fasta object (here called `novelFA`), then for etiquette reasons a valid email address of the person who sends the data to NCBI and the required folder locations on the harddrive. If they are not available, `hoardeR` will create them without asking.

The `xmlFolder` stores the results that are delivered from NCBI in xml format and the `logFolder` stores the log files. These files are especially then important, when the run crashes and should be continued or if the computer is switched off in between.

If a crashed run should be continued, it is enough to run the same command as the initially one using the same fasta object and `hoardeR` will continue from the point, where it crashed or where it was interrupted. If the option `verbose=TRUE` is set (default), `hoardeR` keeps reporting status updates of the blast runs. There are still some fine-tuning parameters available, see the manual for details. However, there are some etiquette

parameters of NCBI that cannot be changed to smaller values than the defaults. These are e.g. the frequency of the requests and the total amount of parallel blast runs.

The verbose output of a `blastSeq` run straight after executing the command looks as follows

```
Missing: 3
Running: 1
Finished: 0
Avg. Blast Time: 00:00:00
Total running time: 00:00:04
-----
```

indicating e.g. that here are still three fasta sequences waiting to be blasted, one sequence is already running an NCBI, zero are finished so far and the running time of the whole run is 4 seconds. After a while (option `delay_rid`), `blastSeq` starts to check, if the active blast runs are finalized. If there are still free blast slots available (maximum number is defined via `n_blast` option), `blastSeq` spawns the next ones. In that case, the verbose output looks like this:

```
Run RW99J31C01R : 00:02:23
Missing: 1
Running: 1
Finished: 2
Avg. Blast Time: 00:01:10
Total running time: 00:02:40
-----
```

Here, we see that there is still one sequence missing and one is also running. The ID number of that run is `RW99J31C01R` and it is active since 2 minutes and 23 seconds. So far, two sequence runs are ready and their average running time was 1 minute and 10 seconds.

The function stops, when there are neither missing nor running sequences anymore. The results of the run are stored in the `xmlFolder`. These results can already be analyzed while the `blastSeq` run is still active, so intermediate results can be obtained on the fly. For that, just open another instance of R and continue to analyze the Blast results.

Analyze the Blast Results

To analyze the blast results, first the xml files have to be imported to `hoardeR`. For that there is the `importXML()` function. It expects the folder address where the xml files are stored, in our example it is the `hoardeROut` folder in the `projFolder` folder.

```
R> xmls <- importXML(folder=file.path(projFolder,"hoardeROut/"))
```

One of the first steps is to check what hit organisms were found. As naturally all sequences are also found in *Bos Taurus*/Cow, the host organism can be excluded from the table

```
R> tableSpecies(xmls, exclude="Bos taurus")
```

This table can be displayed as barplot e.g. in the following way

```
R> par(oma=c(5,0,0,0))
R> barplot(sort(tableSpecies(xmls, exclude="Bos taurus"), decreasing=TRUE), las=2)
```

Next the hits for the *Sus scrofa*/Pig will be visualized. For that, the xml results will be first filtered accordingly. For that also the `tableSpecies` command is used. This time, however, a species is not excluded with the `exclude` option, but explicitly defined with the `species` parameter. Further, instead of just counting the occurrences, we request in addition the location information by setting `locations=TRUE`:

```
R> tableSpecies(xmls, species="Sus scrofa", locations = TRUE)
```

	Organism	hitID	hitLen	hitChr	hitStart	hitEnd	origChr	origStart	origEnd
28	Sus scrofa breed mixed chromosome 4, Sscrofa10.2	494	644	4	105815870	105816509	3	18148822	18149422

without the `locations=TRUE` option, the output is very minimalistic, as it just gives the frequencies:

```
R> tableSpecies(xmls)
```

Bos taurus	Equus caballus	Sus scrofa	Ovis aries
6	1	1	3

Especially when a single species is defined, without setting `locations=TRUE` the function just gives the frequency.

```
R> tableSpecies(xmls, species="Sus scrofa")
```

```
Sus scrofa
1
```

The column `Organism` indicates the hit organism and the corresponding assembly. In this case it is *Sus scrofa* and *Sscrofa10.2*. Cross-checking with the `species` data table reveals, that this assembly is the default assembly at Ensembl:

```
R> species[grepl("Sus scrofa", species$Scientific.name),]
```

	Common.name	Scientific.name	Taxon.ID	Ensembl.Assembly	Accession	Variation.database	Regulation.database	Pre.assembly
57	Pig	Sus scrofa	9823	Sscrofa10.2	GCA_000003025.4	-	Y	-
58	Pig FPC_map (Pre)	Sus scrofa map	NA	-	-	-	-	MAP

That means that the assembly can be obtained automatically from `hoardeR` for further analysis without using any additional parameters. The command `getAnnotation` downloads automatically the corresponding annotation into the specified folder (here: `/home/daniel/annotation`):

```
R> ssannot <- getAnnotation(species = "Sus scrofa",
+                           annotationFolder="/home/daniel/annotation")
```

Having the annotation information available from Ensembl, we can intersect the Blast results with the annotation. First we check if the found loci intersect with an intergenic region in the *Sus scrofa* genome. Here, we assume that more than one hit was found and results would be stored in a list

```
R> pigHits <- tableSpecies(xmls, species="Sus scrofa", locations = TRUE)
R> pigInter <- list()
R> for(i in 1:nrow(pigHits)){
R>   pigInter[[i]] <- intersectXMLAnnot(pigHits[i,], ssannot)
R> }
```

However, unfortunately there is no intersection:

```
R> pigInter
[[1]]
Empty data.table (0 rows) of 15 cols: V1,V2,V3,V4,V5,V6...
```

Hence, we allow for a larger search area and add flanking sites of 100kB to each side of the search area (that means, 100kB are added to each side of the annotated regions), using the option `flanking=100` in the `intersectXMLAnnot()` function:

```
R> pigInter.flank <- list()
R> for(i in 1:nrow(pigHits)){
R>   pigInter.flank[[i]] <- intersectXMLAnnot(pigHits[i,], ssannot, flanking=100)
R> }
# This part transforms the list into a data frame and removes those 'hits' that
# do not report any intersection
```

```
R> pigInter.flank <- pigInter.flank[sapply(pigInter.flank,nrow)>0]
R> pigInter.flank <- do.call(rbind, pigInter.flank)
```

This results in an intersection:

```
R> pigInter.flank
  V1      V2      V3      V4      V5 V6 V7 V8      V9 origChr origStart origEnd hitChr hitStart hitEnd
1:  4 ensembl gene 105858080 105858409 . . . gene_id "ENSSSCG00000006603"; gene_version "2"; gene_source "ensembl"; gene_biotype "protein_coding"; 3 18148822 18149422 4 105815870 105816509
```

The intersections contains first the annotation information as provided by the annotation file and then in addition the location of the hit in the host organism (`origChr`, `origStart`, `origEnd`) and in the hit organism (`hitChr`, `hitStart`, `hitEnd`).

Visualize the intersected Blast results

In order to visualize the results the `plotHit` function can be used. In its basic usage, it only plots the similarity between the original and the hit organism like this:

```
R> plotHit(
+   hits=pigInter.flank,
+   flanking=100,
+   diagonal=0.25 ,
+   hitSpecies = "Sus scrofa",
+   origSpecies = "Bos taurus",
+   fastaFolder = "/home/ejo138/fasta/",
+   # The following options are optional
+   window=NULL ,
+   which=NULL,
+   figureFolder = "/home/daniel/figures/",
+   figurePrefix = "pigIS"
+ )
```

The parameters are then as follows. `hits` expects the `data.frame` from above with the intersected hits. If `hits` receives a `data.frame` with more than one row (i.e. several hits were found) a `figureFolder` should be provided, as for each hit a separate plot will be created. For example, with the above optional parameters a figure for each hit will be created in the folder `/home/daniel/figures/`, using the prefix `pigIS`. However, if only a single figure should be created, without storing it directly to the hard drive, one would drop the `figureFolder` and `figurePrefix` parameters (they are `NULL` by default) and instead specify in the `which` option, which hit should be plotted. An alternative would be to restrict the matrix given to the `hits` option, e.g. for seeing the first hit only, one would use the option `hits=pigInter.flank[1,]`.

The other required options control the behavior of the plot. The `flanking` option defines the plotting area around the hit in Mb and the `diagonal` option is a threshold for the minimum similarity after that a similarity line should be plotted. As higher that value (between 0 and 1) as more restrictive it is and as less similarity lines will be used in the plot. The two parameters `hitSpecies` and `origSpecies` define the two organisms that should be compared. In our case the original organism was *Bos taurus*/cow and the hit was *Sus scrofa*/pig. The location of the corresponding fasta files (or to where they should be stored), is defined in the `fastaFolder` option. Again, if the assembly is the same as in the `species` dataset, the fasta files will be automatically fetched from the Ensembl page. In case a tailored assembly is required, it can be specified with additional options, see the manual or a later example for further instructions on that.

The similarity is calculated using a shifting window approach. That means, the plotting area is divided into chunks of a certain length, defined with the `window` option. By default that window has the similar length as the hit-sequence has, but an own value can be defined in the `window` option. Our function then tests all pairwise combinations between the chunks of the original organism and the hit organism, calculates the similarity between them and stores the best result. That way, each chunk/window of the original organism

gets assigned a chunk from the hit organism. Only those combinations are then considered further, that are at least having a similarity value as defined in `diagonal`.

With that, the figure for the pig would look like this:




Figure 1: Similarity between Pig and Cow

On the top part of the plot the annotation of *Bos taurus* is plotted, on the bottom part the one of *Sus scrofa*, followed by the axes of the chromosomal regions. The vertical center line indicates the hit, further highlighted with the black spots on the chromosomal regions. It appears that this region is rather similar, although in the case of *Bos taurus* it seems that the genomic region is a bit stretched, as the lines are not parallel but open up towards the *Bos taurus* side. Also, it seems that the identified gene in the flanking region corresponds to an already annotated gene in *Bos taurus*.

We continue with examining the same hit for *Ovis aries*/sheep. Here we have

```
R> tableSpecies(xmls, species="Ovis aries", locations = TRUE)
```

	Organism	hitID	hitLen	hitChr	hitStart	hitEnd	origChr	origStart	origEnd
8	Ovis aries breed Texel chromosome 3, Oar_v4.0, whole genome shotgun sequence	577	616	3	33956005	33955391	11	72654673	72655273
16	Ovis aries breed Texel chromosome 14, Oar_v4.0, whole genome shotgun sequence	553	613	14	59259439	59260042	18	62650696	62651296
24	Ovis aries breed Texel chromosome 1, Oar_v4.0, whole genome shotgun sequence	539	604	1	101361981	101362569	3	18148822	18149422

Indicating that the hit comes from the *Oar_v4.0* assembly. However, in Ensembl is only the *Oar_v3.1* assembly and hence, the automatic download would use the wrong assembly and consequently also would operate on the wrong coordinates. Hence, the correct assembly has to be downloaded from the *Ovis aries*

consortium or NCBI and stored as chromosome-wise fasta files in the fasta folder. Same holds also for the annotation files.

From the NCBI page we can download the *Ovis aries* annotation in gff format and we copy the file into the `/home/daniel/Annotations/` folder, using the filename `Ovis_aries.Oar_v4.0.NCBI.chr.gff.gz`. The naming follows again the same scheme, first is the scientific name of the organism with underscores, then is the assembly version, then the release source (or Ensembl release version), followed by the required file ending. Please notice that the annotation needs to be zipped. The annotation can then be loaded to **hoarder** again with `getAnnotation()` command, specifying the required parameters.

```
R> oaannot4.0 <- getAnnotation(species = "Ovis aries", release="NCBI", version="Oar_v4.0",
+                             type="gff", annotationFolder="/home/daniel/Annotations/")
```

The blast hits are imported as before with

```
R> oaHits <- tableSpecies(xmlNew, species="Ovis aries", locations=TRUE)
```

In order to adjust the namings, some additional hand-on is required, e.g. in the following type of way:

```
R> for(i in 0:24){
+   refName <- paste("NC_0",19458+i, ".2", sep="")
+   oaannot4.0$V1[oaannot4.0$V1==refName] <- i+1
+ }
R> i <- 25
R> refName <- paste("NC_0",19458+i, ".2", sep="")
R> oaannot4.0$V1[oaannot4.0$V1==refName] <- "X"
```

With that the intersections between the hits and the annotation using a flanking side of 50Mb can be calculated and then finally merged into a data frame.


```
R> sheepInter4.0 <- list()
R> for(i in 1:nrow(oaHits)){
+   sheepInter4.0[[i]] <- intersectXMLAnnot(oaHits[i,], oaannot4.0, flanking=50)
+ }

R> sheepInter4.0 <- sheepInter4.0[sapply(sheepInter4.0,nrow)>0]
R> sheepInter4.0 <- do.call(rbind, sheepInter4.0)
```

The following command will create then the visualization

```
plotHit(
  hits=sheepInter4.0,
  flanking=50,
  window=NULL,
  diagonal=0.25,
  hitSpecies = "Ovis aries",
  hitSpeciesVersion = "Oar_v4.0",
  origSpecies = "Bos taurus",
  fastaFolder = "/home/daniel/fastas/",
  origAnnot=btannot,
  hitAnnot=oaannot4.0,
  coverage=TRUE,
  bamFolder = "/home/daniel/bams/",
  which=1:2,
  figureFolder = "/home/daniel/figures/",
  figurePrefix = "sheepNOVEL"
)
```

Here, more optional parameters are used. These are, in addition to the already previous used ones, the definition of the used hit species version in `hitSpeciesVersion`. Further, the logical flag `coverage=TRUE` is set. If this is the case, the function expects to find bam files in a folder defined in `bamFolder` that contain information about the coverage. The plot that shows in addition then also the coverage would look this this:



By default all bam files located in the `bamFolder` are taken and the average is taken. However, the filenames of single bam files can also be given to the `bamFiles` option. In case of a case/control study, it might be desirable to plot two different coverage curves. For that, the `groupIndex` defines to which group a bam file belongs and the color of the corresponding coverage is then defined in `groupColor`.